

Lecture 2:

Image Enhancement

What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Image filtering

What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Image filtering

Types of Images

Binary



Types of Images

Binary



Gray Scale



Types of Images

Binary




Gray Scale



Color




Binary image representation




Slide credit: Ulas Bagci

Grayscale image representation



Slide credit: Ulas Bagci

Color Image - one channel



Slide credit: Ulas Bagci



Color image representation



Slide credit: Ulas Bagci

Images are sampled

What happens when we zoom into the images we capture?




Resolution

is a **sampling** parameter, defined in dots per inch (DPI) or equivalent measures of spatial pixel density, and its standard value for recent screen technologies is 72 dpi



Images are Sampled and Quantized

- An image contains discrete number of pixels
 - A simple example
 - Pixel value:
 - “grayscale”
(or “intensity”): $[0, 255]$



Images are Sampled and Quantized

- An image contains discrete number of pixels

- A simple example

- Pixel value:

- “grayscale”


(or “intensity”): $[0, 255]$

- “color”

- RGB: $[R, G, B]$

- Lab: $[L, a, b]$

- HSV: $[H, S, V]$




What we will learn today?



- Image sampling and quantization
- Image histograms
- Images as functions
- Image filtering

Histogram

- Histogram of an image provides the frequency of the brightness (intensity) value in the image.
- Histogram captures the distribution of gray levels in the image.
- How frequently each gray level occurs in the image





Histogram



Slide credit: Dr. Mubarak Shah

Histogram – Use case



Slide credit: Dr. Mubarak Shah


What we will learn today?

- Image sampling and quantization
- Image histograms
- **Images as functions**
- Image filtering

Images as discrete functions

- Images are usually **digital (discrete)**:
 - **Sample** the 2D space on a regular grid
- Represented as a matrix of integer values

pixel



j	62	79	23	119	120	05	4	0
i	10	10	9	62	12	78	34	0
	10	58	197	46	46	0	0	48
	176	135	5	188	191	68	0	49
	2	1	1	29	26	37	0	77
	0	89	144	147	187	102	62	208
	255	252	0	166	123	62	0	31
	166	63	127	17	1	0	99	30

Images as coordinates

Cartesian coordinates

$$I[n, m] = \begin{bmatrix} \ddots & & & & \\ & I[-1, 1] & I[0, 1] & I[1, 1] & \\ \dots & I[-1, 0] & \underline{I[0, 0]} & I[1, 0] & \dots \\ & I[-1, -1] & I[0, -1] & I[1, -1] & \\ & & \vdots & & \ddots \end{bmatrix}$$

$I[n, m]$ is the notation for discrete functions


Images as functions

- **An Image** as a function f from \mathbb{R}^2 to \mathbb{R}^M :
 - $I(x, y)$ gives the **intensity** at position (x, y)
 - Defined over a rectangle, with a finite range:

$$I: [a,b] \times [c,d] \rightarrow [0,255]$$

Domain
support

range



Images as functions

- **An Image** as a function f from \mathbb{R}^2 to \mathbb{R}^M :
 - $I(x, y)$ gives the **intensity** at position (x, y)
 - Defined over a rectangle, with a finite range:

$$I: \underbrace{[a, b] \times [c, d]}_{\substack{\text{Domain} \\ \text{support}}} \rightarrow \underbrace{[0, 255]}_{\text{range}}$$

- A color image: $I(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$

What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Image filtering

Systems and Filters

Filtering:

- Forming a new image whose pixel values are transformed from original pixel values

Goals:

- Goal is to extract useful information from images, or transform images into another domain where we can modify/enhance image properties
 - Features (edges, corners, blobs...)
 - super-resolution; in-painting; de-noising


De-noising



Salt and pepper noise



Super-resolution



In-painting



Bertamio et al

Image filtering

- Image filters in spatial domain
 - Filter is a mathematical operation of a grid of numbers
 - Smoothing, sharpening, measuring texture
- Image filters in the frequency domain
 - Filtering is a way to modify the frequencies of images
 - Denoising, sampling, image compression

Image filtering

- Image filtering:
 - Compute function of local neighborhood at each position

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

Image filtering

- Image filtering:
 - Compute function of local neighborhood at each position

`h=output f=filter I=image`

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k,n+l]$$

`2d coords=k,l 2d coords=m,n`

`[] [] []`

Example: box filter

$$f[\cdot, \cdot]$$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Image filtering

AI Academy

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot] \frac{1}{9}$$


$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

Image filtering



1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot] \frac{1}{9}$$

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10							

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

Image filtering

AI Academy

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot] \frac{1}{9}$$

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20						

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

Image filtering

AI Academy

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot] \frac{1}{9}$$

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30					

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

Image filtering

AI Academy

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot] \frac{1}{9}$$

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30				

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

1	1	1
1	1	1
1	1	1

1	1	1
1	1	1
1	1	1

$$h[.,.]$$
[illegible]

Credit: S. Seitz

Image filtering

AI Academy

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot] \frac{1}{9}$$

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30				
						?			
				50					

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

Image filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$$\frac{1}{9} f[\cdot, \cdot]$$

1	1	1
1	1	1
1	1	1

Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)
- Why does it sum to one?

$$\frac{1}{9} f[\cdot, \cdot]$$

1	1	1
1	1	1
1	1	1

Smoothing with box filter




Image filtering

- Image filtering:
 - Compute function of local neighborhood at each position

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

- Really important!
 - Enhance images
 - Denoise, resize, increase contrast, etc.
 - Extract information from images
 - Texture, edges, distinctive points, etc.
 - Detect patterns
 - Template matching

Think-Pair-Share time



1.

0	0	0
0	1	0
0	0	0

2.

0	0	0
0	0	1
0	0	0

3.

1	0	-1
2	0	-2
1	0	-1

4.

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

1. Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

?

1. Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

2. Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

?

2. Practice with linear filters




Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

3. Practice with linear filters




1	0	-1
2	0	-2
1	0	-1

Sobel




Vertical Edge
(absolute value)

3. Practice with linear filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)

4. Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)

4. Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$


1	1	1
1	1	1
1	1	1




Sharpening filter

- Accentuates differences with local average

4. Practice with linear filters



before



after

Correlation and Convolution

- 2d correlation

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k,n+l]$$

`h=filter2(f,I);` or `h=imfilter(I,f);`

- 2d convolution

$$h[m,n] = \sum_{k,l} f[k,l] I[m-k,n-l]$$



`h=conv2(f,I);` or `h=imfilter(I,f,'conv');`

Correlation and convolution are identical when the filter is symmetric.

1D Discrete convolution (symbol:) *

We are going to convolve a function **f** with a filter **h**.

$$g[n] = \sum_k f[k]h[n - k]$$






1D Discrete convolution (symbol: $*$)

We are going to convolve a function **f** with a filter **h**.

$$g[n] = \sum_k f[k]h[n - k]$$




We first need to calculate $h[n-k, m-l]$



Discrete convolution (symbol: $*$)

$*$




We are going to convolve a function **f** with a filter **h**.



Discrete convolution (symbol: $*$)

$*$




We are going to convolve a function **f** with a filter **h**.



Discrete convolution (symbol: $*$)

 $*$



We are going to convolve a function **f** with a filter **h**.



Discrete convolution (symbol:)

*



We are going to convolve a function **f** with a filter **h**.



Discrete convolution (symbol: $*$)

$*$


We are going to convolve a function **f** with a filter **h**.



Discrete convolution (symbol: $*$)

$*$

We are going to convolve a function **f** with a filter **h**.




2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.


n

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter(h[,]) that is 3x3. and an image (f[,]) that is 7x7.


n

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.


n

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.


n

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter(h[,]) that is 3x3. and an image (f[,]) that is 7x7.


n

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter(h[,]) that is 3x3. and an image (f[,]) that is 7x7.

2D convolution example

1	2	3
4	5	6
7	8	9

Input

		m		
n		-1	0	1
	-1	-1	-2	-1
	0	0	0	0
	1	1	2	1

Kernel

-13	-20	-17
-18	-24	-18
13	20	17

Output

2D convolution example

1	2	1	
0	0	0	3
-1	-2	-1	6
	7	8	9

$$\begin{aligned} &= x[-1,-1] \cdot h[1,1] + x[0,-1] \cdot h[0,1] + x[1,-1] \cdot h[-1,1] \\ &\quad + x[-1,0] \cdot h[1,0] + x[0,0] \cdot h[0,0] + x[1,0] \cdot h[-1,0] \\ &\quad + x[-1,1] \cdot h[1,-1] + x[0,1] \cdot h[0,-1] + x[1,1] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 0 \cdot 0 + 1 \cdot 0 + 2 \cdot 0 + 0 \cdot (-1) + 4 \cdot (-2) + 5 \cdot (-1) = -13 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

2D convolution example

1	2	1
0	0	0
1	2	3
-1	-2	-1
4	5	6
7	8	9

$$\begin{aligned} &= x[0,-1] \cdot h[1,1] + x[1,-1] \cdot h[0,1] + x[2,-1] \cdot h[-1,1] \\ &\quad + x[0,0] \cdot h[1,0] + x[1,0] \cdot h[0,0] + x[2,0] \cdot h[-1,0] \\ &\quad + x[0,1] \cdot h[1,-1] + x[1,1] \cdot h[0,-1] + x[2,1] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 1 \cdot 0 + 2 \cdot 0 + 3 \cdot 0 + 4 \cdot (-1) + 5 \cdot (-2) + 6 \cdot (-1) = -20 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

2D convolution example

		1	2	1
1	0	2	0	0
4	-1	5	-2	-1
7	8	9		

$$\begin{aligned}
 &= x[1,-1] \cdot h[1,1] + x[2,-1] \cdot h[0,1] + x[3,-1] \cdot h[-1,1] \\
 &\quad + x[1,0] \cdot h[1,0] + x[2,0] \cdot h[0,0] + x[3,0] \cdot h[-1,0] \\
 &\quad + x[1,1] \cdot h[1,-1] + x[2,1] \cdot h[0,-1] + x[3,1] \cdot h[-1,-1] \\
 &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 2 \cdot 0 + 3 \cdot 0 + 0 \cdot 0 + 5 \cdot (-1) + 6 \cdot (-2) + 0 \cdot (-1) = -17
 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

2D convolution example

1	2	1	
	1	2	3
0	0	0	
	4	5	6
-1	-2	-1	
	7	8	9

$$\begin{aligned} &= x[-1,0] \cdot h[1,1] + x[0,0] \cdot h[0,1] + x[1,0] \cdot h[-1,1] \\ &\quad + x[-1,1] \cdot h[1,0] + x[0,1] \cdot h[0,0] + x[1,1] \cdot h[-1,0] \\ &\quad + x[-1,2] \cdot h[1,-1] + x[0,2] \cdot h[0,-1] + x[1,2] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 1 \cdot 2 + 2 \cdot 1 + 0 \cdot 0 + 4 \cdot 0 + 5 \cdot 0 + 0 \cdot (-1) + 7 \cdot (-2) + 8 \cdot (-1) = -18 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

2D convolution example

1 1	2 2	1 3
0 4	0 5	0 6
-1 7	-2 8	-1 9

$$\begin{aligned} &= x[0,0] \cdot h[1,1] + x[1,0] \cdot h[0,1] + x[2,0] \cdot h[-1,1] \\ &\quad + x[0,1] \cdot h[1,0] + x[1,1] \cdot h[0,0] + x[2,1] \cdot h[-1,0] \\ &\quad + x[0,2] \cdot h[1,-1] + x[1,2] \cdot h[0,-1] + x[2,2] \cdot h[-1,-1] \\ &= 1 \cdot 1 + 2 \cdot 2 + 3 \cdot 1 + 4 \cdot 0 + 5 \cdot 0 + 6 \cdot 0 + 7 \cdot (-1) + 8 \cdot (-2) + 9 \cdot (-1) = -24 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

2D convolution example

1	2	3	
4	5	6	
7	8	9	

$$\begin{aligned} &= x[1,0] \cdot h[1,1] + x[2,0] \cdot h[0,1] + x[3,0] \cdot h[-1,1] \\ &\quad + x[1,1] \cdot h[1,0] + x[2,1] \cdot h[0,0] + x[3,1] \cdot h[-1,0] \\ &\quad + x[1,2] \cdot h[1,-1] + x[2,2] \cdot h[0,-1] + x[3,2] \cdot h[-1,-1] \\ &= 2 \cdot 1 + 3 \cdot 2 + 0 \cdot 1 + 5 \cdot 0 + 6 \cdot 0 + 0 \cdot 0 + 8 \cdot (-1) + 9 \cdot (-2) + 0 \cdot (-1) = -18 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Convolution in 2D - examples



Original

*

•0	•0	•0
•0	•1	•0
•0	•0	•0

=

?

Convolution in 2D - examples



Original



•0	•0	•0
•0	•1	•0
•0	•0	•0



Filtered
(no change)

Convolution in 2D - examples



Original

*

•0	•0	•0
•0	•0	•1
•0	•0	•0

=

?

Convolution in 2D - examples



Original



•0	•0	•0
•0	•0	•1
•0	•0	•0



Shifted right
By 1 pixel

Convolution in 2D - examples



Original

$$* \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array} = ?$$

Convolution in 2D - examples



Original

*

$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1

=



Blur (with a
box filter)

Convolution in 2D - examples



Original

•0	•0	•0
•0	•2	•0
•0	•0	•0

—

$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1

= ?

(Note that filter sums to 1)

“details of the image”

•0	•0	•0
•0	•1	•0
•0	•0	•0

+

•0	•0	•0
•0	•1	•0
•0	•0	•0

—


$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1


- What does blurring take away?



—




=



- Let's add it back:



+ a



=



Convolution in 2D – Sharpening filter



Original

•0	•0	•0
•0	•2	•0
•0	•0	•0

–

$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1

=



Sharpening filter: Accentuates differences with local average

Convolution properties

Commutative: $a * b = b * a$

- Conceptually no difference between filter and signal
- But particular filtering implementations might break this equality, e.g., image edges

Associative: $a * (b * c) = (a * b) * c$

- Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
- This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$

Convolution properties

Commutative: $a * b = b * a$

- Conceptually no difference between filter and signal
- But particular filtering implementations might break this equality, e.g., image edges


Associative: $a * (b * c) = (a * b) * c$

- Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
- This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Correlation is not associative
- Why important?


Convolution properties

- Commutative: $a * b = b * a$
 - Conceptually no difference between filter and signal
 - But particular filtering implementations might break this equality, e.g., image edges
- Associative: $a * (b * c) = (a * b) * c$
 - Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
 - Correlation is not associative (rotation effect)
 - Why important?
- Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out: $ka * b = a * kb = k(a * b)$

Smoothing with Gaussian filter



Smoothing with box filter



Gaussian filters

- Remove “high-frequency” components from the image (low-pass filter)
 - Images become more smooth
- Gaussian convolved with Gaussian...
 - ...is another Gaussian
 - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
 - Convoluting two times with Gaussian kernel of width σ is same as convoluting once with kernel of width $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into product of two 1D Gaussians

Separability of the Gaussian filter

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp -\frac{x^2 + y^2}{2\sigma^2} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp -\frac{x^2}{2\sigma^2} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp -\frac{y^2}{2\sigma^2} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability example

2D convolution
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors
into a product of 1D
filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution
along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution
along the remaining column:

Separability

Why is separability useful in practice?

Separability

Why is separability useful in practice?

$M \times N$ image, $P \times Q$ filter

- 2D convolution: $\sim MN PQ$ multiply-adds
- Separable 2D: $\sim MN(P+Q)$ multiply-adds

Speed up = $PQ/(P+Q)$

9x9 filter = $\sim 4.5x$ faster


Practical matters

How big should the filter be?

- Values at edges should be near zero
- Gaussians have infinite extent...
- Rule of thumb for Gaussian: set filter half-width to about 3σ

Practical matters

- What about near the edge?
 - the filter window falls off the edge of the image
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge



Convolution in Convolutional Neural Networks

- Convolution is the basic operation in CNNs
- Learning convolution kernels allows us to learn which `features' provide useful information in images.

NON-LINEAR FILTERS

Median filters

- Operates over a window by selecting the median intensity in the window.
- ‘Rank’ filter as based on ordering of gray levels
 - E.G., min, max, range filters


1	1	1
1	1	1
1	1	1

1	1	1
1	1	1
1	1	1

$$h[.,.]$$
[illegible]

Credit: S. Seitz

Image filtering - mean



$$f[\cdot, \cdot] \frac{1}{9}$$

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30				
				50					

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

Median filter?

$I[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

				?					

Median filters

- Operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?

Noisy Jack – Salt and Pepper



Mean Jack – 3 x 3 filter



Very Mean Jack – 11 x 11 filter



Noisy Jack – Salt and Pepper



Median Jack – 3 x 3



Very Median Jack – 11 x 11



Median filters

- Operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?
- Is a median filter a kind of convolution?

Review: questions

1. Write down a 3x3 filter that both:
 - Returns a positive value if the average value of the 4-adjacent neighbors is less than the center,
 - Returns a negative value otherwise.

Review: questions

1. Write down a 3x3 filter that both:

- Returns a positive value if the average value of the 4-adjacent neighbors is less than the center,
 - Returns a negative value otherwise.
- $$\begin{bmatrix} 0 & -1/4 & 0 \\ -1/4 & 1 & -1/4 \\ 0 & -1/4 & 0 \end{bmatrix}$$